

Predicting Software Support Ticket Criticality

Using Sparse Imbalanced Data and NLP

April 14, 2020

Matt Kane
Nick Sylva
Elena Petrov

Problem Statement

Create a screening tool for incoming email support requests that predicts the severity of the issue to facilitate support ticket triage before a human can read the support email.

Dataset


*Real business dataset spanning
6,000 clients, 425,000 tickets,
and ~2.7 million ticket actions
over the last 10 years*

- Sourced from Greenshades, a payroll and employee services software provider for SMBs
- Raw data was ~4.3GB in CSV format


Dataset Considerations

- Large dataset
- Thorough data cleansing
- Semantic augmentation
- Multi-label classification problem with highly imbalanced dataset


Histogram of Note Lengths (No Truncation)
Baseline Dataset 1




Histogram of Note Lengths (No Truncation)
Baseline Dataset 2




Histogram of Note Lengths (Truncated)
Baseline Dataset 1




Histogram of Note Lengths (Truncated)
Baseline Dataset 2



Baseline 1 Preprocessed Raw Dataset
Ticket Priority as Percentage of Total



Baseline 2 01_preprocessed
Ticket Priority as Percentage of Total



Model Specifications

Model 1: Simple Classifier with BERT Fine Tuning

- Simplest model
- BERT to single Dense layer
- Sequence representation
- Examine effectiveness of fine-tuning

Model 2: Frozen BERT and Dense NN



- Frozen BERT: no fine-tuning
- Sequence representation
- 4 dense layers
- Examine if dense network can pick up features better than BERT fine-tuning alone

Model 3: Frozen BERT, CNN, and Dense NN



- Frozen BERT: no fine-tuning
- Full embeddings for each token
- Passed through three filter sizes of 64 filters each to pooling, concatenation, and dense network
- Pick up on relationships between token embeddings

Model 1 Results: Simple Classifier with BERT Fine Tuning

Baseline 1: Regular Normalization




Baseline 2: Semantic Augmentation










Model 2: Frozen BERT and Dense NN

Baseline 1: Regular Normalization




Baseline 2: Semantic Augmentation




Model 3 Results: Frozen BERT, CNN, and Dense NN

Baseline 1: Regular Normalization


Confusion Matrix
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 1 Dataset



Confusion Matrix
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 1 Downsampled Dataset




Confusion Matrix
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 1 SMOTE Dataset




Baseline 2: Semantic Augmentation


Confusion Matrix
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 2 Dataset




Confusion Matrix
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 2 Downsampled Dataset




Confusion Matrix
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 2 SMOTE Dataset




Receiver Operating Characteristic for
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 1 Dataset




Receiver Operating Characteristic for
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 1 Downsampled Dataset




Receiver Operating Characteristic for
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 1 SMOTE Dataset




Receiver Operating Characteristic for
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 2 Dataset



Receiver Operating Characteristic for
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 2 Downsampled Dataset



Receiver Operating Characteristic for
BERT Frozen Weights with FF CNN to a Dense NN Using
Baseline 2 SMOTE Dataset



Conclusions

- Fine-tuning on BERT does not solve everything
 - Bespoke dataset that is riddled with colloquialisms, abbreviations, sentence fragments and domain-specific language
- Other potential enhancements:
 - Use unfrozen BERT weights with model 3 architecture
 - Hyperparameter tuning

Results


Model	Dataset	Balancing	Accuracy
Model 3	Baseline 1	None	84.89%
Model 2	Baseline 1	None	81.11%
Model 2	Baseline 2	None	79.51%
Model 1	Baseline 2	SMOTE	78.55%
Model 1	Baseline 2	None	78.53%
Model 1	Baseline 1	None	78.16%
Model 1	Baseline 1	Downsampled	78.09%
Model 3	Baseline 1	SMOTE	65.76%
Model 2	Baseline 1	SMOTE	61.20%
Model 3	Baseline 2	SMOTE	60.71%
Model 2	Baseline 2	SMOTE	58.05%
Model 3	Baseline 1	Downsampled	55.90%
Model 2	Baseline 1	Downsampled	52.69%
Model 3	Baseline 2	Downsampled	52.06%
Model 2	Baseline 2	Downsampled	47.53%
Model 1	Baseline 2	Downsampled	16.72%
Model 1	Baseline 1	SMOTE	4.56%
Model 3	Baseline 2	*** In Process ***	

Notes:

Model 1: BERT Fine Tuning Model

Model 2: BERT Frozen Weights with FF Dense NN

Model 3: BERT BERT Frozen Weights with FF CNN to a Dense NN



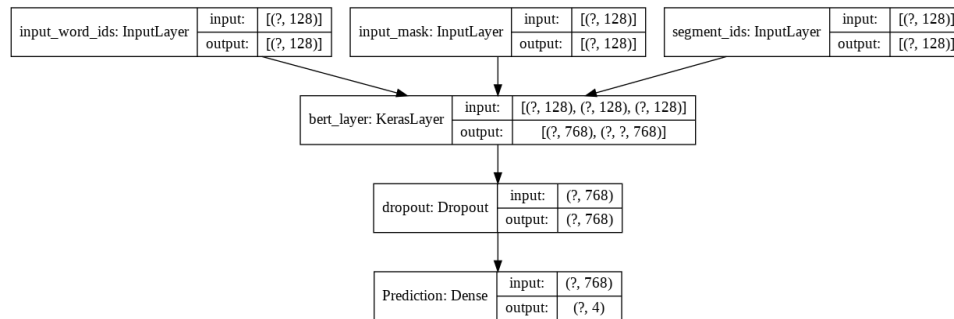
Questions



Appendices

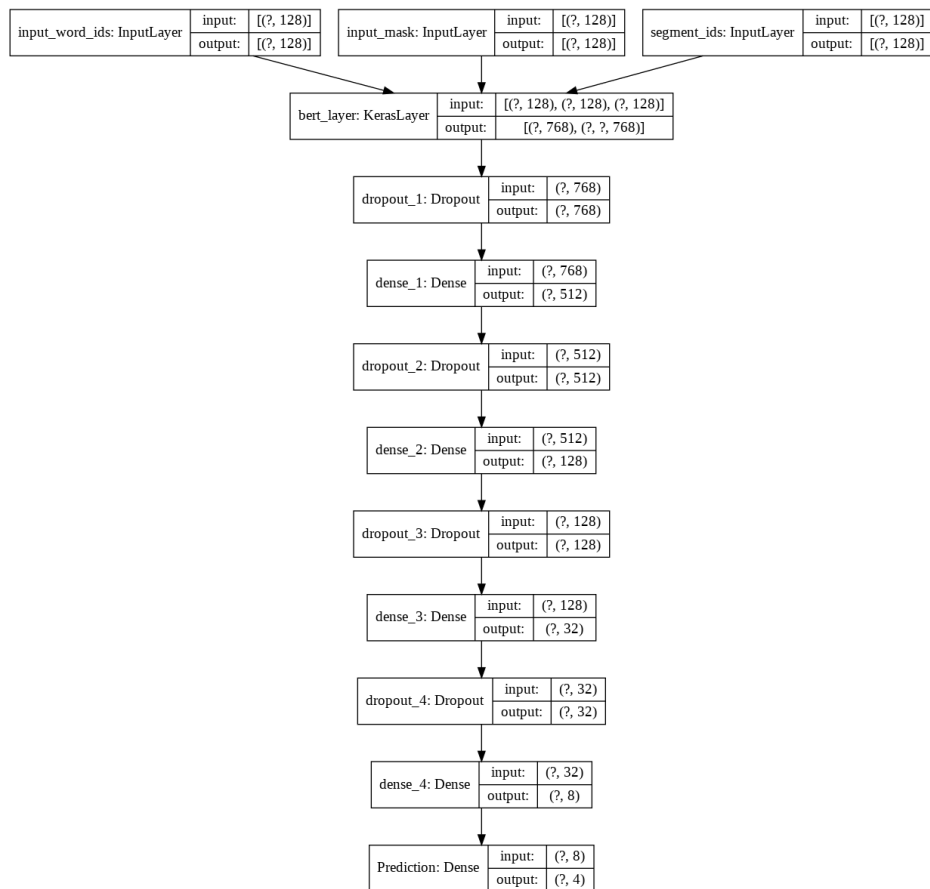
Model 1 Diagram

*Simple Classifier with BERT
Fine Tuning*



Model 2 Diagram

*Frozen BERT and Dense
Neural Network*



Model 3 Diagram

*Frozen BERT, CNN and
Feed Forward Dense
Neural Network*

